AI/ML Engineer – LLM Applications

Company: Northwind Analytics (fictional)
Location: Bengaluru, India (Hybrid – 3 days in office)
Experience: 3–6 years
Employment type: Full-time

About the role
We are building an AI assistant that answers questions over enterprise documents. You will design, build and ship LLM-powered features end to end, from retrieval pipelines to production APIs.

What you will do
- Build retrieval-augmented generation (RAG) pipelines over large document collections
- Design prompts, evaluation sets and guardrails for LLM features
- Develop and maintain Python APIs (FastAPI) that serve models in production
- Work with vector databases (Pinecone, FAISS or similar) and embeddings
- Deploy and monitor services on AWS using Docker
- Collaborate with product and data teams to measure quality and cost

Must have
- 3+ years of professional experience in Python
- Hands-on experience with LLMs (OpenAI, Claude, Gemini or open-source models)
- RAG and LangChain or LlamaIndex
- Vector databases and embeddings
- FastAPI or Flask for REST APIs
- Docker and AWS

- SQL

Nice to have
- Fine-tuning (LoRA / PEFT), Hugging Face Transformers
- LLM evaluation frameworks (RAGAS, DeepEval)
- Kubernetes, CI/CD

Education
B.E./B.Tech/M.Tech in Computer Science or a related field.